

House Price Prediction Model

Import Libraries and Helpers

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

Load and Inspect Dataset

Loading House Prediction Dataset

```
df = pd.read_csv("train.csv")
df.head()

   Id MSSubClass MSZoning LotFrontage LotArea Street Alley LotShape \
0   1          60      RL     65.0    8450    Pave   NaN    Reg
1   2          20      RL     80.0    9600    Pave   NaN    Reg
2   3          60      RL     68.0   11250    Pave   NaN   IR1
3   4          70      RL     60.0    9550    Pave   NaN   IR1
4   5          60      RL     84.0   14260    Pave   NaN   IR1

   LandContour Utilities ... PoolArea PoolQC Fence MiscFeature MiscVal MoSold \
0       Lvl    AllPub ...      0    NaN    Nan      NaN      0      2
1       Lvl    AllPub ...      0    NaN    Nan      NaN      0      5
2       Lvl    AllPub ...      0    NaN    Nan      NaN      0      9
3       Lvl    AllPub ...      0    NaN    Nan      NaN      0      2
```

```

4      Lvl    AllPub ...      0   NaN   NaN      NaN      0     12
YrSold SaleType SaleCondition SalePrice
0    2008       WD      Normal  208500
1    2007       WD      Normal  181500
2    2008       WD      Normal  223500
3    2006       WD  Abnorml  140000
4    2008       WD      Normal  250000
[5 rows x 81 columns]

```

Inspecting Dataset

```

# checking data types
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 81 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Id               1460 non-null   int64  
 1   MSSubClass        1460 non-null   int64  
 2   MSZoning          1460 non-null   object  
 3   LotFrontage       1201 non-null   float64 
 4   LotArea            1460 non-null   int64  
 5   Street             1460 non-null   object  
 6   Alley              91 non-null    object  
 7   LotShape            1460 non-null   object  
 8   LandContour         1460 non-null   object  
 9   Utilities           1460 non-null   object  
 10  LotConfig           1460 non-null   object  
 11  LandSlope           1460 non-null   object  
 12  Neighborhood        1460 non-null   object  
 13  Condition1          1460 non-null   object  
 14  Condition2          1460 non-null   object  
 15  BldgType            1460 non-null   object  
 16  HouseStyle          1460 non-null   object  
 17  OverallQual         1460 non-null   int64  
 18  OverallCond         1460 non-null   int64  
 19  YearBuilt            1460 non-null   int64  
 20  YearRemodAdd        1460 non-null   int64  
 21  RoofStyle            1460 non-null   object  
 22  RoofMatl             1460 non-null   object  
 23  Exterior1st          1460 non-null   object  
 24  Exterior2nd          1460 non-null   object  
 25  MasVnrType           588 non-null    object  
 26  MasVnrArea           1452 non-null   float64 
 27  ExterQual            1460 non-null   object  
 28  ExterCond            1460 non-null   object 

```

29	Foundation	1460	non-null	object
30	BsmtQual	1423	non-null	object
31	BsmtCond	1423	non-null	object
32	BsmtExposure	1422	non-null	object
33	BsmtFinType1	1423	non-null	object
34	BsmtFinSF1	1460	non-null	int64
35	BsmtFinType2	1422	non-null	object
36	BsmtFinSF2	1460	non-null	int64
37	BsmtUnfSF	1460	non-null	int64
38	TotalBsmtSF	1460	non-null	int64
39	Heating	1460	non-null	object
40	HeatingQC	1460	non-null	object
41	CentralAir	1460	non-null	object
42	Electrical	1459	non-null	object
43	1stFlrSF	1460	non-null	int64
44	2ndFlrSF	1460	non-null	int64
45	LowQualFinSF	1460	non-null	int64
46	GrLivArea	1460	non-null	int64
47	BsmtFullBath	1460	non-null	int64
48	BsmtHalfBath	1460	non-null	int64
49	FullBath	1460	non-null	int64
50	HalfBath	1460	non-null	int64
51	BedroomAbvGr	1460	non-null	int64
52	KitchenAbvGr	1460	non-null	int64
53	KitchenQual	1460	non-null	object
54	TotRmsAbvGrd	1460	non-null	int64
55	Functional	1460	non-null	object
56	Fireplaces	1460	non-null	int64
57	FireplaceQu	770	non-null	object
58	GarageType	1379	non-null	object
59	GarageYrBlt	1379	non-null	float64
60	GarageFinish	1379	non-null	object
61	GarageCars	1460	non-null	int64
62	GarageArea	1460	non-null	int64
63	GarageQual	1379	non-null	object
64	GarageCond	1379	non-null	object
65	PavedDrive	1460	non-null	object
66	WoodDeckSF	1460	non-null	int64
67	OpenPorchSF	1460	non-null	int64
68	EnclosedPorch	1460	non-null	int64
69	3SsnPorch	1460	non-null	int64
70	ScreenPorch	1460	non-null	int64
71	PoolArea	1460	non-null	int64
72	PoolQC	7	non-null	object
73	Fence	281	non-null	object
74	MiscFeature	54	non-null	object
75	MiscVal	1460	non-null	int64
76	MoSold	1460	non-null	int64
77	YrSold	1460	non-null	int64

```
78 SaleType      1460 non-null  object
79 SaleCondition 1460 non-null  object
80 SalePrice     1460 non-null  int64
dtypes: float64(3), int64(35), object(43)
memory usage: 924.0+ KB
```

```
# summary statistics
df.describe()
```

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	\
count	1460.000000	1460.000000	1201.000000	1460.000000	1460.000000	
mean	730.500000	56.897260	70.049958	10516.828082	6.099315	
std	421.610009	42.300571	24.284752	9981.264932	1.382997	
min	1.000000	20.000000	21.000000	1300.000000	1.000000	
25%	365.750000	20.000000	59.000000	7553.500000	5.000000	
50%	730.500000	50.000000	69.000000	9478.500000	6.000000	
75%	1095.250000	70.000000	80.000000	11601.500000	7.000000	
max	1460.000000	190.000000	313.000000	215245.000000	10.000000	

	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea	BsmtFinSF1	...	\
count	1460.000000	1460.000000	1460.000000	1452.000000	1460.000000	...	
mean	5.575342	1971.267808	1984.865753	103.685262	443.639726	...	
std	1.112799	30.202904	20.645407	181.066207	456.098091	...	
min	1.000000	1872.000000	1950.000000	0.000000	0.000000	...	
25%	5.000000	1954.000000	1967.000000	0.000000	0.000000	...	
50%	5.000000	1973.000000	1994.000000	0.000000	383.500000	...	
75%	6.000000	2000.000000	2004.000000	166.000000	712.250000	...	
max	9.000000	2010.000000	2010.000000	1600.000000	5644.000000	...	

	WoodDeckSF	OpenPorchSF	EnclosedPorch	3SsnPorch	ScreenPorch	\
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	
mean	94.244521	46.660274	21.954110	3.409589	15.060959	
std	125.338794	66.256028	61.119149	29.317331	55.757415	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	25.000000	0.000000	0.000000	0.000000	
75%	168.000000	68.000000	0.000000	0.000000	0.000000	
max	857.000000	547.000000	552.000000	508.000000	480.000000	

	PoolArea	MiscVal	MoSold	YrSold	SalePrice
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000
mean	2.758904	43.489041	6.321918	2007.815753	180921.195890
std	40.177307	496.123024	2.703626	1.328095	79442.502883
min	0.000000	0.000000	1.000000	2006.000000	34900.000000
25%	0.000000	0.000000	5.000000	2007.000000	129975.000000
50%	0.000000	0.000000	6.000000	2008.000000	163000.000000
75%	0.000000	0.000000	8.000000	2009.000000	214000.000000
max	738.000000	15500.000000	12.000000	2010.000000	755000.000000

```
[8 rows x 38 columns]
```

```
# checking for missing values
df.isna().sum().sort_values(ascending=False)

PoolQC      1453
MiscFeature 1406
Alley       1369
Fence       1179
MasVnrType  872
...
MoSold        0
YrSold        0
SaleType       0
SaleCondition  0
SalePrice      0
Length: 81, dtype: int64
```

Visual and Insights

```
# identifying numeric and categorical columns
numeric_features = df.select_dtypes(include=["int64", "float64"]).columns.tolist()
categorical_features = df.select_dtypes(include=["object"]).columns.tolist()

numeric_features, categorical_features[:10] # preview first 10 categories

(['Id',
 'MSSubClass',
 'LotFrontage',
 'LotArea',
 'OverallQual',
 'OverallCond',
 'YearBuilt',
 'YearRemodAdd',
 'MasVnrArea',
 'BsmtFinSF1',
 'BsmtFinSF2',
 'BsmtUnfSF',
 'TotalBsmtSF',
 '1stFlrSF',
 '2ndFlrSF',
 'LowQualFinSF',
 'GrLivArea',
 'BsmtFullBath',
 'BsmtHalfBath',
 'FullBath',
 'HalfBath',
 'BedroomAbvGr',
 'KitchenAbvGr',
 'TotRmsAbvGrd',
 'Fireplaces',
 'GarageYrBlt',
```


```
'GarageCars',
'GarageArea',
'WoodDeckSF',
'OpenPorchSF',
'EnclosedPorch',
'3SsnPorch',
'ScreenPorch',
'PoolArea',
'MiscVal',
'MoSold',
'YrSold',
'SalePrice'],
['MSZoning',
'Street',
'Alley',
'LotShape',
'LandContour',
'Utilities',
'LotConfig',
'LandSlope',
'Neighborhood',
'Condition1'])
```

```
sns.set(style="whitegrid")
```

SalePrice Distribution - Histogram

House prices often follow a right-skewed distribution, meaning most homes are moderately priced, with fewer extremely expensive properties.

```
plt.figure(figsize=(10,5))
sns.histplot(df["SalePrice"], kde=True, color="skyblue")
plt.title("SalePrice Distribution")
plt.xlabel("SalePrice")
plt.ylabel("Count")
plt.show()
```



Insight:


SalePrice is strongly right-skewed. A log transformation reduces skew and stabilizes variance, which often improves regression model performance.

Numeric Features vs SalePrice


I have visualized here how major numeric predictors relate to SalePrice. I have used these 3 most important features:

- GrLivArea — Above ground living area
- TotalBsmtSF — Basement area
- OverallQual — Quality rating (1–10)


```
plt.figure(figsize=(8,5))
sns.scatterplot(data=df, x="GrLivArea", y="SalePrice", alpha=0.6)
plt.title("GrLivArea vs SalePrice")
plt.show()
```



```
plt.figure(figsize=(8,5))
sns.scatterplot(data=df, x="TotalBsmtSF", y="SalePrice", alpha=0.6, color="orange")
plt.title("TotalBsmtSF vs SalePrice")
plt.show()
```



```
plt.figure(figsize=(8,5))
sns.scatterplot(data=df, x="OverallQual", y="SalePrice", alpha=0.6, color="green")
plt.title("OverallQual vs SalePrice")
plt.show()
```



Insight:


- GrLivArea and TotalBsmtSF show strong positive relationships with SalePrice — larger homes sell for more.
- OverallQual has a very strong upward trend: better quality homes command significantly higher prices.

Categorical Features vs SalePrice


Using boxplot here to compare price ranges across categorical groups. I have used these 2 most common features:

- Neighborhood
- HouseStyle

```
# boxplot 1 - Neighborhood vs SalePrice
plt.figure(figsize=(14,6))
sns.boxplot(data=df, x="Neighborhood", y="SalePrice")
plt.xticks(rotation=90)
plt.title("Neighborhood vs SalePrice")
plt.show()
```



```
# boxplot 2 - HouseStyle vs SalePrice
plt.figure(figsize=(10,5))
sns.boxplot(data=df, x="HouseStyle", y="SalePrice")
plt.title("HouseStyle vs SalePrice")
plt.show()
```




Insight:

Different neighborhoods show significantly different price levels, indicating strong location effects. House-style types also show variation, with certain styles typically commanding higher prices.

Correlation Heatmap

```
plt.figure(figsize=(18,10))
corr = df[numerical_features].corr()
sns.heatmap(corr, annot=False, cmap="coolwarm")
plt.title("Correlation Heatmap (Numeric Features)")
plt.show()
```



Insight:

The strongest correlations with SalePrice include:

- OverallQual
- GrLivArea
- GarageCars & GarageArea
- TotalBsmtSF

These features will likely be highly important for prediction.

Preprocessing and Split Dataset

In this section:

- Remove non-informative columns
- Separate features (X) and target (y)
- Optionally apply a log transform to SalePrice
- Identify numeric and categorical features
- Build a preprocessing pipeline with imputation, scaling and one-hot encoding
- Split the data into training and test sets

Feature and Target Selection

```
# Drop unnecessary non-features
df = df.drop(columns=["Id"])

# Target
y = df["SalePrice"]
# Features
X = df.drop(columns=["SalePrice"])
```

House prices are very skewed; models work better on log prices.

```
# Log-transform target to reduce skew
y_log = np.log1p(y)    # log(1 + Price)

y_original = y          # keep for reference
y = y_log               # from now on use y as log-price
```

Train-Test Split Dataset

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

X_train.shape, X_test.shape
```

```
((1168, 79), (292, 79))
```

Identify numeric & categorical feature lists (on X, not df)

```
numeric_features = X.select_dtypes(include=["int64", "float64"]).columns.tolist()
categorical_features = X.select_dtypes(include=["object"]).columns.tolist()

numeric_features[:10], categorical_features[:10] # quick preview

(['MSSubClass',
 'LotFrontage',
 'LotArea',
 'OverallQual',
 'OverallCond',
 'YearBuilt',
 'YearRemodAdd',
 'MasVnrArea',
 'BsmtFinSF1',
 'BsmtFinSF2'],
 ['MSZoning',
 'Street',
 'Alley',
 'LotShape',
 'LandContour',
 'Utilities',
 'LotConfig',
 'LandSlope',
 'Neighborhood',
 'Condition1'])
```

Building Pipelines for Feature(X) Preprocessing

Numeric pipeline: impute (median) + scale

This pipeline will fillup the missing numeric values with median and scale them.

```
numeric_transformer = Pipeline([
    ("imputer", SimpleImputer(strategy="median")),
    ("scaler", StandardScaler())
])
```

Categorical pipeline: impute (most frequent) + one-hot encode

This pipeline will fillup the missing categorical values with the most frequent ones and scale with one-hot encode

```
categorical_transformer = Pipeline(steps=[
    ("imputer", SimpleImputer(strategy="most_frequent")),
    ("onehot", OneHotEncoder(handle_unknown="ignore"))
])
```

Combining the numerical and categorical pipelines with ColumnTransformer

```
preprocessor = ColumnTransformer(  
    transformers=[  
        ("num", numeric_transformer, numeric_features),  
        ("cat", categorical_transformer, categorical_features),  
    ]  
)
```

Baseline Linear Regression Model Train and Evaluation

Building the pipeline - Preprocessor + LinearRegression

```
# baseline linear regression pipeline  
baseline_linreg = Pipeline([  
    ("preprocessor", preprocessor),  
    ("linreg", LinearRegression())  
)  
  
baseline_linreg  
  
Pipeline(steps=[('preprocessor',  
    ColumnTransformer(transformers=[('num',  
        Pipeline(steps=[('imputer',  
            SimpleImputer(strategy='median')),  
            ('scaler',  
                StandardScaler()))]),  
        ['MSSubClass', 'LotFrontage',  
        'LotArea', 'OverallQual',  
        'OverallCond', 'YearBuilt',  
        'YearRemodAdd', 'MasVnrArea',  
        'BsmtFinSF1', 'BsmtFinSF2',  
        'BsmtUnfSF', 'TotalBsmtSF',  
        '1stFlrSF', '2ndFlrSF',  
        'LowQual...',  
        'LandContour', 'Utilities',  
        'LotConfig', 'LandSlope',  
        'Neighborhood', 'Condition1',  
        'Condition2', 'BldgType',  
        'HouseStyle', 'RoofStyle',  
        'RoofMatl', 'Exterior1st',  
        'Exterior2nd', 'MasVnrType',  
        'ExterQual', 'ExterCond',  
        'Foundation', 'BsmtQual',  
        'BsmtCond', 'BsmtExposure',  
        'BsmtFinType1',  
        'BsmtFinType2', 'Heating',  
        'HeatingQC', 'CentralAir',
```

```
        'Electrical', ...]])),  
    ('linreg', LinearRegression()))])
```

Fitting the model on Training Data

```
baseline_linreg.fit(X_train, y_train)  
  
Pipeline(steps=[('preprocessor',  
                 ColumnTransformer(transformers=[('num',  
                                              Pipeline(steps=[('imputer',  
                                                 SimpleImputer(strategy='median')),  
                                              ('scaler',  
                                               StandardScaler())]),  
                                              ['MSSubClass', 'LotFrontage',  
                                               'LotArea', 'OverallQual',  
                                               'OverallCond', 'YearBuilt',  
                                               'YearRemodAdd', 'MasVnrArea',  
                                               'BsmtFinSF1', 'BsmtFinSF2',  
                                               'BsmtUnfSF', 'TotalBsmtSF',  
                                               '1stFlrSF', '2ndFlrSF',  
                                               'LowQual...',  
                                               'LandContour', 'Utilities',  
                                               'LotConfig', 'LandSlope',  
                                               'Neighborhood', 'Condition1',  
                                               'Condition2', 'BldgType',  
                                               'HouseStyle', 'RoofStyle',  
                                               'RoofMatl', 'Exterior1st',  
                                               'Exterior2nd', 'MasVnrType',  
                                               'ExterQual', 'ExterCond',  
                                               'Foundation', 'BsmtQual',  
                                               'BsmtCond', 'BsmtExposure',  
                                               'BsmtFinType1',  
                                               'BsmtFinType2', 'Heating',  
                                               'HeatingQC', 'CentralAir',  
                                               'Electrical', ...]])),  
    ('linreg', LinearRegression()))])
```

Predicting on Test Data

```
y_pred = baseline_linreg.predict(X_test)
```

Evaluating Baseline Linear Regression Model (y in log-scale)

```
# evaluating baseline linear regression  
mse_log = mean_squared_error(y_test, y_pred)  
rmse_log = np.sqrt(mse_log)  
mae_log = mean_absolute_error(y_test, y_pred)  
r2_log = r2_score(y_test, y_pred)  
  
print(f"RMSE : {rmse_log:.4f}")
```

```
print(f"MAE : {mae_log:.4f}")
print(f"R2 : {r2_log:.4f}")
```

```
RMSE : 0.1281
MAE  : 0.0883
R2   : 0.9121
```

Evaluating Baseline Linear Regression Model (y in original-scale)

```
# converting predictions back to original SalePrice scale
y_test_orig = np.expm1(y_test)
y_pred_orig = np.expm1(y_pred)

mse_orig = mean_squared_error(y_test_orig, y_pred_orig)
rmse_orig = np.sqrt(mse_orig)
mae_orig = mean_absolute_error(y_test_orig, y_pred_orig)
r2_orig = r2_score(y_test_orig, y_pred_orig)

print(f"RMSE (original): {rmse_orig:.2f}$")
print(f"MAE  (original): {mae_orig:.2f}$")
print(f"R2  (original): {r2_orig:.4f}$")

RMSE (original): 22,739.77$
MAE  (original): 14,899.28$
R2  (original): 0.9326$
```

Train and Evaluate with other Models + Comparison

Defining Other Models

```
models = {}

# 1. Linear Regression(untuned)
models["Linear Regression"] = Pipeline([
    ("preprocessor", preprocessor),
    ("linreg", LinearRegression())
])

# 2. Ridge Regression
models["Ridge Regression"] = Pipeline([
    ("preprocessor", preprocessor),
    ("ridgereg", Ridge())
])

# 3. Lasso Regression
models["Lasso Regression"] = Pipeline([
    ("preprocessor", preprocessor),
    ("lassoreg", Lasso())
])
```

```

# 4. Random Forest Regressor
models["Random Forest Regressor"] = Pipeline([
    ("preprocessor", preprocessor),
    ("rfreg", RandomForestRegressor(
        n_estimators=100,
        random_state=42,
        n_jobs=-1
    ))
])

# 5. Gradient Boosting Regressor
models["Gradient Boosting Regressor"] = Pipeline([
    ("preprocessor", preprocessor),
    ("gradreg", GradientBoostingRegressor(random_state=42))
])

```

Defining Tuned Models(Random Forest and Gradient Boosting) using GridSearchCV

```

# tuning random forest
rf_pipeline = Pipeline([
    ("preprocessor", preprocessor),
    ("rfreg", RandomForestRegressor(random_state=42, n_jobs=-1))
])

rf_param_grid = {
    "rfreg__n_estimators": [100, 200],
    "rfreg__max_depth": [None, 10, 20],
    "rfreg__max_features": ["sqrt", "log2"],
    "rfreg__min_samples_split": [2, 5]
}

rf_grid_search = GridSearchCV(
    estimator=rf_pipeline,
    param_grid=rf_param_grid,
    cv=5,
    scoring="neg_mean_squared_error",
    n_jobs=-1,
    verbose=1
)

rf_grid_search.fit(X_train, y_train)

print("Best parameters (Random Forest):")
print(rf_grid_search.best_params_)

best_rf_rmse_cv = np.sqrt(-rf_grid_search.best_score_)
print(f"Best CV RMSE (log-space): {best_rf_rmse_cv:.4f}")

best_rf_pipeline = rf_grid_search.best_estimator_
models["Random Forest Regressor (Tuned)"] = best_rf_pipeline

```

```

Fitting 5 folds for each of 24 candidates, totalling 120 fits
Best parameters (Random Forest):
{'rfreg__max_depth': None, 'rfreg__max_features': 'sqrt', 'rfreg__min_samples_split': 2, 'rfreg__n_estimators': 200}
Best CV RMSE (log-space): 0.1479

# tuning gradient boosting regressor
gb_pipeline = Pipeline([
    ("preprocessor", preprocessor),
    ("gradreg", GradientBoostingRegressor(random_state=42))
])

# hyperparameter grid for Gradient Boosting
gb_param_grid = {
    "gradreg__n_estimators": [100, 200],
    "gradreg__learning_rate": [0.05, 0.1, 0.2],
    "gradreg__max_depth": [2, 3, 4]
}

gb_grid_search = GridSearchCV(
    estimator=gb_pipeline,
    param_grid=gb_param_grid,
    cv=5,
    scoring="neg_mean_squared_error",
    n_jobs=-1,
    verbose=1
)

gb_grid_search.fit(X_train, y_train)

print("Best parameters (Gradient Boosting):")
print(gb_grid_search.best_params_)

best_gb_rmse_cv = np.sqrt(-gb_grid_search.best_score_)
print(f"Best CV RMSE (log-space): {best_gb_rmse_cv:.4f}")

best_gb_pipeline = gb_grid_search.best_estimator_
models["Gradient Boosting Regressor (Tuned)"] = best_gb_pipeline

Fitting 5 folds for each of 18 candidates, totalling 90 fits
Best parameters (Gradient Boosting):
{'gradreg__learning_rate': 0.05, 'gradreg__max_depth': 3, 'gradreg__n_estimators': 200}
Best CV RMSE (log-space): 0.1327

```

Evaluating All Models

```

# helper function to evaluate with all models
def model_evaluation(name, model, X_train, y_train, X_test, y_test):
    # train model
    model.fit(X_train, y_train)
    # predictions(in log space)

```

```

y_pred = model.predict(X_test)

# convert back to original sale prices
y_pred_orig = np.expm1(y_pred)
y_test_orig = np.expm1(y_test)

# metrics
mse = mean_squared_error(y_test_orig, y_pred_orig)
rmse = np.sqrt(mse)
mae = mean_absolute_error(y_test_orig, y_pred_orig)
r2 = r2_score(y_test_orig, y_pred_orig)

print(f"{name:35s} -> RMSE : {rmse:.4f} | MAE : {mae:.2f} | R² : {r2:.4f}")

return rmse, mae, r2

# evaluating all models and storing results
results = []

for name, model in models.items():
    rmse, mae, r2 = model_evaluation(name, model, X_train, y_train, X_test, y_test)
    results.append({
        "Model": name,
        "RMSE": rmse,
        "MAE": mae,
        "R²": r2
    })

Linear Regression      -> RMSE : 22739.7741 | MAE : 14,899.28 | R² : 0.9326
Ridge Regression       -> RMSE : 23856.6945 | MAE : 15,738.72 | R² : 0.9258
Lasso Regression       -> RMSE : 88270.8537 | MAE : 59,931.24 | R² : -0.0158
Random Forest Regressor -> RMSE : 29963.8159 | MAE : 17,712.67 | R² : 0.8829
Gradient Boosting Regressor -> RMSE : 30019.7658 | MAE : 16,926.73 | R² : 0.8825
Random Forest Regressor (Tuned) -> RMSE : 34765.3740 | MAE : 18,482.06 | R² : 0.8424
Gradient Boosting Regressor (Tuned) -> RMSE : 30761.7621 | MAE : 16,742.79 | R² : 0.8766

# comparison dataframe with all model's result
results_df = pd.DataFrame(results)
results_df = results_df.sort_values(by="RMSE", ascending=True).reset_index(drop=True)
results_df

```

	Model	RMSE	MAE	R ²
0	Linear Regression	22739.774051	14899.276389	0.932585
1	Ridge Regression	23856.694526	15738.719034	0.925800
2	Random Forest Regressor	29963.815911	17712.671661	0.882948
3	Gradient Boosting Regressor	30019.765785	16926.734665	0.882510
4	Gradient Boosting Regressor (Tuned)	30761.762091	16742.794622	0.876630
5	Random Forest Regressor (Tuned)	34765.373965	18482.059355	0.842428
6	Lasso Regression	88270.853655	59931.238025	-0.015829

Final Model and Interpretation

Choosing the best model based on lowest RMSE

```
# selecting the best model based on lowest RMSE
best_row = results_df.sort_values("RMSE").iloc[0]
best_model_name = best_row["Model"]

print("Best model based on RMSE:", best_model_name)
print(best_row)

Best model based on RMSE: Linear Regression
Model      Linear Regression
RMSE      22739.774051
MAE       14899.276389
R²        0.932585
Name: 0, dtype: object

# retrieving the best model pipeline from the models dictionary
best_model = models[best_model_name]
best_model

Pipeline(steps=[('preprocessor',
                 ColumnTransformer(transformers=[('num',
                                                   Pipeline(steps=[('imputer',
                                                               SimpleImputer(strategy='median'))),
                                                   ('scaler',
                                                   StandardScaler()))]),
                 ['MSSubClass', 'LotFrontage',
                  'LotArea', 'OverallQual',
                  'OverallCond', 'YearBuilt',
                  'YearRemodAdd', 'MasVnrArea',
                  'BsmtFinSF1', 'BsmtFinSF2',
                  'BsmtUnfSF', 'TotalBsmtSF',
                  '1stFlrSF', '2ndFlrSF',
                  'LowQual...',
                  'LandContour', 'Utilities',
                  'LotConfig', 'LandSlope',
                  'Neighborhood', 'Condition1',
                  'Condition2', 'BldgType',
                  'HouseStyle', 'RoofStyle',
                  'RoofMatl', 'Exterior1st',
                  'Exterior2nd', 'MasVnrType',
                  'ExterQual', 'ExterCond',
                  'Foundation', 'BsmtQual',
                  'BsmtCond', 'BsmtExposure',
                  'BsmtFinType1',
                  'BsmtFinType2', 'Heating',
                  'HeatingQC', 'CentralAir',
```


```
        'Electrical', ...]])),  
    ('linreg', LinearRegression()))])
```

Refitting and Evaluation with Best Model

```
# training the best model on the training data  
best_model.fit(X_train, y_train)  
  
# predictions with the best model (log SalePrice)  
y_pred_log = best_model.predict(X_test)  
  
# converting back to original SalePrice scale  
y_test_orig = np.expm1(y_test)  
y_pred_orig = np.expm1(y_pred_log)  
  
# final evaluation metrics  
final_mse = mean_squared_error(y_test_orig, y_pred_orig)  
final_rmse = np.sqrt(final_mse)  
final_mae = mean_absolute_error(y_test_orig, y_pred_orig)  
final_r2 = r2_score(y_test_orig, y_pred_orig)  
  
print(f"Final Model: {best_model_name}")  
print(f"RMSE (original $): {final_rmse:.2f}")  
print(f"MAE (original $): {final_mae:.2f}")  
print(f"R2 : {final_r2:.4f}")  
  
Final Model: Linear Regression  
RMSE (original $): 22,739.77  
MAE (original $): 14,899.28  
R2 : 0.9326
```

Plotting Predicted vs Actual SalePrice

```
# predicted vs actual prices  
plt.figure(figsize=(7,6))  
sns.scatterplot(x=y_test_orig, y=y_pred_orig, alpha=0.6)  
max_price = max(y_test_orig.max(), y_pred_orig.max())  
min_price = min(y_test_orig.min(), y_pred_orig.min())  
  
plt.plot([min_price, max_price], [min_price, max_price], "r--") # ideal line  
plt.xlabel("Actual SalePrice")  
plt.ylabel("Predicted SalePrice")  
plt.title(f"Predicted vs Actual SalePrice ({best_model_name})")  
plt.show()
```



Feature Importance

```
# extracting feature names after preprocessing
feature_names = best_model.named_steps["preprocessor"].get_feature_names_out()


# get the final estimator inside the pipeline (last step)
final_step_name = list(best_model.named_steps.keys())[-1]
estimator = best_model.named_steps[final_step_name]

# get importance values
if hasattr(estimator, "feature_importances_"):
    importances = estimator.feature_importances_
elif hasattr(estimator, "coef_"):
    # for linear models, use absolute value of coefficients
    importances = np.abs(estimator.coef_)
else:
    importances = None
    print("This estimator does not provide feature_importances_ or coef_.")
```

```
if importances is not None:
    # build a DataFrame of feature importances
    fi_df = pd.DataFrame({
        "Feature": feature_names,
        "Importance": importances
    }).sort_values(by="Importance", ascending=False)

    # show top 15 features
    fi_top15 = fi_df.head(15)
    fi_top15

# plotting top features
if importances is not None:
    plt.figure(figsize=(8,6))
    sns.barplot(data=fi_top15, x="Importance", y="Feature")
    plt.title(f"Top 15 Features - {best_model_name}")
    plt.tight_layout()
    plt.show()
```



Final Model Summary

The best-performing model based on lowest RMSE was **Linear Regression**. Although more complex models were tested (Ridge, Lasso, Random Forest, Gradient Boosting), Linear Regression provided the strongest generalization performance on the test set.

Performance (Original SalePrice Scale)

- RMSE: \$22,739.77
- MAE: \$14,899.28
- R²: 0.9326

These results indicate that the model explains approximately 93% of the variation in house prices, with average prediction errors around \$15,000. The predicted vs actual plot shows a strong linear relationship, and residuals are centered around zero, suggesting no major bias.

Interpretation

Key predictors include property quality (OverallQual), above-ground living area (GrLivArea), garage capacity, and basement size. These variables strongly influence housing prices, which aligns with common real-estate valuation factors.

Conclusion

Linear Regression, combined with proper preprocessing (imputation, scaling, encoding), provides an effective and interpretable baseline for house price prediction. The model performs well across diverse homes and can be further improved through additional feature engineering or more advanced regression techniques.